

In [68]:

```
import pandas as pd
import numpy as np
import warnings
```

In [69]:

```
df_train = pd.read_csv('../datasets/house_train.csv')
```

In [70]:

```
df_train
```

Out[70]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	M
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	
...
1455	1456	60	RL	62.0	7917	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	
1456	1457	20	RL	85.0	13175	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	MnPrv	NaN	
1457	1458	70	RL	66.0	9042	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	GdPrv	Shed	
1458	1459	20	RL	68.0	9717	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	
1459	1460	20	RL	75.0	9937	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	

1460 rows × 81 columns

In [71]:

```
df_train.columns
```

Out[71]:

```
Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',
       'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',
       'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType',
       'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd',
       'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',
       'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual',
       'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',
       'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating',
       'HeatingQC', 'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF',
       'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath',
       'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual',
       'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType',
       'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual',
       'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF',
       'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'PoolQC',
       'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',
       'SaleCondition', 'SalePrice'],
      dtype='object')
```

In [72]:

```
df_train.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Id                  1460 non-null  int64
1   MSSubClass          1460 non-null  int64
2   MSZoning            1460 non-null  object
3   LotFrontage        1201 non-null  float64
4   LotArea            1460 non-null  int64
5   Street             1460 non-null  object
6   Alley              91 non-null    object
7   LotShape           1460 non-null  object
8   LandContour        1460 non-null  object
9   Utilities          1460 non-null  object
10  LotConfig          1460 non-null  object
11  LandSlope          1460 non-null  object
12  Neighborhood        1460 non-null  object
13  Condition1         1460 non-null  object
14  Condition2         1460 non-null  object
```

In [73]:

```
df_train['SalePrice'].hist()
```

Out[73]:

<AxesSubplot:>




In [74]:

```
plt.hist(df_train['SalePrice'])
```

Out[74]:

```
(array([148., 723., 373., 135., 51., 19., 4., 3., 2., 2.]),  
array([ 34900., 106910., 178920., 250930., 322940., 394950., 466960.,  
        538970., 610980., 682990., 755000.]),  
<BarContainer object of 10 artists>)
```



In [75]:

```
df_train['SalePrice'].describe()
```

Out[75]:


```
count      1460.000000  
mean      180921.195890  
std       79442.502883  
min       34900.000000  
25%      129975.000000  
50%      163000.000000  
75%      214000.000000  
max       755000.000000  
Name: SalePrice, dtype: float64
```

In [76]:

```
plt.boxplot(df_train['SalePrice'])
```

Out[76]:

```
{'whiskers': [matplotlib.lines.Line2D at 0x134d6a9a880>,
<matplotlib.lines.Line2D at 0x134d6a9abe0>],
'caps': [matplotlib.lines.Line2D at 0x134d6a9af40>,
<matplotlib.lines.Line2D at 0x134d6aa42e0>],
'boxes': [matplotlib.lines.Line2D at 0x134d6a9a520>],
'medians': [matplotlib.lines.Line2D at 0x134d6aa4670>],
'fliers': [matplotlib.lines.Line2D at 0x134d6aa49d0>],
'means': []}
```



In [77]:

```
z_scores = (df_train['SalePrice']-df_train['SalePrice'].mean())/df_train['SalePrice'].std()
df_train = df_train[((z_scores > -3) & (z_scores <= 3))]
```

In [78]:

```
z_scores
```

Out[78]:

```
0      0.347154
1      0.007286
2      0.535970
3     -0.515105
4      0.869545
...
1455  -0.074534
1456   0.366036
1457   1.077242
1458  -0.488356
1459  -0.420697
Name: SalePrice, Length: 1460, dtype: float64
```

In [79]:

```
df_train.head()
```

Out[79]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	0
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	0
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0

5 rows × 81 columns

In [80]:

```
df_train['Alley'].value_counts()
```

Out[80]:

```
Grv1    50
Pave    41
Name: Alley, dtype: int64
```

In [81]:

```
list1=['Alley', 'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2', 'FireplaceQu', 'GarageType',
'GarageFinish', 'GarageQual', 'GarageCond', 'Fence', 'MiscFeature', 'PoolQC']
```

In [82]:

```
dict1 = {}
for i in df_train.columns.drop('SalePrice'):
    if i not in list1:
        dict1[i] = ['', '#N/A', '#N/A N/A', '#NA', '-1.#IND', '-1.#QNAN', '-NaN', '-nan',
                    '1.#IND', '1.#QNAN', '<NA>', 'N/A', 'NA', 'NULL', 'NaN', 'n/a', 'nan', 'null']
    else:
        dict1[i] = ['', '#N/A', '#N/A N/A', '#NA', '-1.#IND', '-1.#QNAN', '-NaN', '-nan',
                    '1.#IND', '1.#QNAN', '<NA>', 'N/A', 'NULL', 'NaN', 'n/a', 'nan', 'null']
len(dict1)
```

Out[82]:

80

In [83]:

```
df_train = pd.read_csv('../datasets/house_train.csv', keep_default_na=False, na_values=dict1)
df_test = pd.read_csv('../datasets/house_test.csv', keep_default_na=False, na_values=dict1)
```

In [84]:

```
df_train.isna().sum()[df_train.isna().sum()>0]
```

Out[84]:

```
LotFrontage    259
MasVnrType      8
MasVnrArea      8
Electrical      1
GarageYrBlt    81
dtype: int64
```

In [85]:

```
df_test.isna().sum()[df_test.isna().sum()>0]
```

Out[85]:

```
MSZoning      4
LotFrontage   227
Utilities      2
Exterior1st    1
Exterior2nd    1
MasVnrType     16
MasVnrArea     15
BsmtFinSF1      1
BsmtFinSF2      1
BsmtUnfSF       1
TotalBsmtSF     1
BsmtFullBath    2
BsmtHalfBath    2
KitchenQual     1
Functional      2
GarageYrBlt     78
GarageCars      1
GarageArea      1
SaleType        1
dtype: int64
```

In [86]:

```
columns_cat = list(df_train.select_dtypes(include='object').columns)
columns_num = list(df_train.select_dtypes(exclude='object').columns)
columns_num.remove('Id')
columns_num.remove('SalePrice')
```

In [87]:

```
from sklearn.impute import SimpleImputer
```

In [88]:

```
imputer_num = SimpleImputer(strategy='median')
imputer_cat = SimpleImputer(strategy='most_frequent')

imputer_num.fit(df_train[columns_num])
imputer_cat.fit(df_train[columns_cat])

df_train[columns_num] = imputer_num.transform(df_train[columns_num])
df_train[columns_cat] = imputer_cat.transform(df_train[columns_cat])

df_test[columns_num] = imputer_num.transform(df_test[columns_num])
df_test[columns_cat] = imputer_cat.transform(df_test[columns_cat])
```

In [89]:

```
df_train.isna().sum()[df_train.isna().sum()>0]
```

Out[89]:

Series([], dtype: int64)

In [90]:

```
df_test.isna().sum()[df_test.isna().sum()>0]
```

Out[90]:

Series([], dtype: int64)

In [91]:

```
from sklearn.preprocessing import MinMaxScaler
```

In [92]:

```
scaler = MinMaxScaler()
scaler.fit(df_train[columns_num])
df_train[columns_num] = scaler.transform(df_train[columns_num])
df_test[columns_num] = scaler.transform(df_test[columns_num])
```

In [93]:

```
from sklearn.preprocessing import OneHotEncoder
```

In [94]:

```
ohe = OneHotEncoder(handle_unknown='ignore')
ohe.fit(df_train[columns_cat])
df_train[ohe.get_feature_names()] = ohe.transform(df_train[columns_cat]).toarray()
df_test[ohe.get_feature_names()] = ohe.transform(df_test[columns_cat]).toarray()
```

In [95]:

```
df_train[ohe.get_feature_names()]
```

Out[95]:

	x0_C (all)	x0_FV	x0_RH	x0_RL	x0_RM	x1_Grvt	x1_Pave	x2_Grvt	x2_NA	x2_Pave	...	x41_ConLw	x41_New	x41_Oth	x41_WD	x42_Abnorml	x42_
0	0.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	0.0	...	0.0	0.0	0.0	1.0	0.0	
1	0.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	0.0	...	0.0	0.0	0.0	1.0	0.0	
2	0.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	0.0	...	0.0	0.0	0.0	1.0	0.0	
3	0.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	0.0	...	0.0	0.0	0.0	1.0	1.0	
4	0.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	0.0	...	0.0	0.0	0.0	1.0	0.0	
...	
1455	0.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	0.0	...	0.0	0.0	0.0	1.0	0.0	
1456	0.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	0.0	...	0.0	0.0	0.0	1.0	0.0	
1457	0.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	0.0	...	0.0	0.0	0.0	1.0	0.0	
1458	0.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	0.0	...	0.0	0.0	0.0	1.0	0.0	
1459	0.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	0.0	...	0.0	0.0	0.0	1.0	0.0	

1460 rows × 266 columns

In [96]:

```
len(ohe.get_feature_names())
```

Out[96]:

266

In [97]:

```
ohe.transform(df_train[columns_cat]).toarray()
```

Out[97]:

```
array([[0., 0., 0., ..., 0., 1., 0.],
       [0., 0., 0., ..., 0., 1., 0.],
       [0., 0., 0., ..., 0., 1., 0.],
       ...,
       [0., 0., 0., ..., 0., 1., 0.],
       [0., 0., 0., ..., 0., 1., 0.],
       [0., 0., 0., ..., 0., 1., 0.]])
```

In [98]:

```
df_train.shape, df_test.shape
```

Out[98]:

```
((1460, 347), (1459, 346))
```

In [99]:

```
column_list = columns_num + list(ohe.get_feature_names()) + ['SalePrice']

corr_values = df_train[column_list].corr()['SalePrice']
corr_values

selected_col = list((corr_values[(corr_values > 0.1) | (corr_values < -0.1)]).index)
selected_col.remove('SalePrice')
```

In [100]:

```
X = df_train[selected_col]
y = df_train['SalePrice']
```

In [101]:

```
from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(X,y)
yp = model.predict(df_test[selected_col])
df_test['SalePrice'] = yp
df_test[['Id', 'SalePrice']].to_csv('finalhpl.csv', index = False)
model.score(X,y)
```

Out[101]:

```
0.8886568944585275
```

In [102]:

```
from sklearn.neighbors import KNeighborsRegressor
modelKNN = KNeighborsRegressor(n_neighbors=3)
modelKNN.fit(df_train[columns_num + list(ohe.get_feature_names())], df_train['SalePrice'])
print(modelKNN.score(df_train[columns_num + list(ohe.get_feature_names())], df_train['SalePrice']))
yp = modelKNN.predict(df_test[columns_num + list(ohe.get_feature_names())])
df_test['SalePrice'] = yp
df_test[['Id', 'SalePrice']].to_csv('subKNN3.csv', index=False)
```

```
0.8672683240814711
```

In [104]:

```
from sklearn.ensemble import RandomForestRegressor

params = {'n_estimators':[20,30,40,50], 'max_depth':[2,3,4,5,6], 'min_samples_leaf':[3,4,5,6]}
gridCV = GridSearchCV(RandomForestRegressor(random_state=51),
                      param_grid=params, cv=5, verbose=0, scoring='accuracy')

gridCV.fit(df_train[selected_col], df_train['SalePrice'])
```

```
ValueError: Classification metrics can't handle a mix of multiclass and continuous targets
```

```
warnings.warn(
C:\Users\Rishabh jain\anaconda3\lib\site-packages\sklearn\model_selection\_validation.py:683: UserWarning: Scoring failed. The score on this train-test partition for these parameters will be set to nan. Details:
Traceback (most recent call last):
  File "C:\Users\Rishabh jain\anaconda3\lib\site-packages\sklearn\model_selection\_validation.py", line 674, in _score
    scores = scorer(estimator, X_test, y_test)
  File "C:\Users\Rishabh jain\anaconda3\lib\site-packages\sklearn\metrics\_scorer.py", line 199, in __call__
    return self._score(partial(_cached_call, None), estimator, X, y_true,
  File "C:\Users\Rishabh jain\anaconda3\lib\site-packages\sklearn\metrics\_scorer.py", line 242, in _score
    return self._sign * self._score_func(y_true, y_pred,
  File "C:\Users\Rishabh jain\anaconda3\lib\site-packages\sklearn\utils\_validation.py", line 63, in inner_f
    return f(*args, **kwargs)
  File "C:\Users\Rishabh jain\anaconda3\lib\site-packages\sklearn\metrics\_classification.py", line 202, in accuracy_score
    y_type, y_true, y_pred = _check_targets(y_true, y_pred)
  File "C:\Users\Rishabh jain\anaconda3\lib\site-packages\sklearn\metrics\_classification.py", line 92, in _check_target
    raise ValueError("Classification metrics can't handle a mix of {0} "
```

In [105]:

```
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
gridCV.best_params_

modelRF = RandomForestRegressor(max_depth=12, min_samples_leaf=2, random_state=95)

modelRF.fit(df_train[selected_col], df_train['SalePrice'])
print(modelRF.score(df_train[selected_col], df_train['SalePrice']))

yp = modelRF.predict(df_test[selected_col])

df_test['SalePrice'] = yp
df_test[['Id', 'SalePrice']].to_csv('modelRF3.csv', index=False)
```

0.9724212494478934

In [106]:

```
from xgboost import XGBRegressor
modelXG = XGBRegressor()
modelXG.fit(df_train[selected_col], df_train['SalePrice'])

yp = modelXG.predict(df_test[selected_col])

df_test['SalePrice'] = yp
df_test[['Id', 'SalePrice']].to_csv('modelXG2.csv', index=False)

modelXG.score(df_train[selected_col], df_train['SalePrice'])
```

Out[106]:

0.9997100838688698

In []:

In []:

In []:

In []:

In [48]:

In [49]:

In [50]:

In []: